

A large, stylized graphic of a sphere composed of many overlapping, concentric circles in shades of purple and lavender, creating a complex, geometric pattern.

watsonx.governance™

watsonx.governance Hands-On Workshop

*Introduction to AI Governance via **Building, Deploying, and Managing** AI models.*

Created by: **Haneen Bakbak**, AI/MLOps Technical Specialist
Marianna Speranza, Brand Technical Specialist
(Haneen.bakbak@ibm.com, Marianna.Speranza@ibm.com),
IBM Canada Data Platform Technical Sales

For more information, visit us at: <https://www.ibm.com/products/watsonx-governance>

Notices and Disclaimers

This guide contains proprietary information which is protected by copyright. No part of this document may be photocopied, reproduced, or translated into another language without a legal license agreement from IBM Corporation.

© Copyright IBM Corporation 2025.

The information contained in these materials is provided for informational purposes only and is provided AS IS without warranty of any kind, express or implied. IBM shall not be responsible for any damages arising out of the use of, or otherwise related to, these materials. Nothing contained in these materials is intended to, nor shall have the effect of, creating any warranties or representations from IBM or its suppliers or licensors, or altering the terms and conditions of the applicable license agreement governing the use of IBM software. References in these materials to IBM products, programs, or services do not imply that they will be available in all countries in which IBM operates. This information is based on current IBM product plans and strategy, which are subject to change by IBM without notice. Product release dates and/or capabilities referenced in these materials may change at any time at IBM's sole discretion based on market opportunities or other factors and are not intended to be a commitment to future product or feature availability in any way.

This document is current as of the initial date of publication and may be changed by IBM at any time. Not all offerings are available in every country in which IBM operates.

IBM, the IBM logo and ibm.com are trademarks or registered trademarks of International Business Machines Corporation in the United States, other countries, or both. If these and other IBM trademarked terms are marked on their first occurrence in this information with a trademark symbol (® or ™), these symbols indicate U.S. registered or common law trademarks owned by IBM at the time this information was published. Such trademarks may also be registered or common law trademarks in other countries. A current list of IBM trademarks is available on the Web at "Copyright and trademark information" at ibm.com/legal/copytrade.shtml

Other company, product and service names may be trademarks or service marks of others.

Table of Contents

Introduction to watsonx.governance	4
Create an Inventory	5
Create an AI Use Case	6
Risk Assessment	11
EU AI Act applicability Assessment	16
Submit for stakeholder review	18
Create a Prompt Template	22
Model Parameters	27
Saving Prompt	28
Evaluating your prompt template	30
Evaluation Results	32
AI Factsheet	33
Govern your prompt template	37
Deployment Space	40
Tracking Prompt Template in a Use Case	39
Update your prompt template	44
Evaluation Studio	46
Drift Monitoring	59
Governance Console	66
Appendix: Navigating the UI of the Governance Console	70
Get Started with your watsonx.governance Workshop	70
watsonx.governance User Interface	71
Changing Colour Theme	71
The Primary Menu	73
Search Use Cases	75
Create Search Filters	77
Configuring your home page: create a new panel	79
Navigation from an object	83
Audit Trail	86
Optional: Promoting your prompt to Validation	86
Deploying the prompt	90
Activate prompt monitoring (Evaluation)	92

Introduction to watsonx.governance

watsonx.governance is a Modern AI Governance toolkit that provides a sleek and intuitive User Interface. Whether Users are looking for responsible AI workflows, or to reduce costs and comply with regulations, watsonx.governance provides the capabilities to empower everyone in the organization with the insights needed to positively impact governance.

This workshop is designed to give you an opportunity to discover the ease of use of watsonx.governance, IBM's AI Governance solution that enhances the efficiency and capabilities of AI developers, users, policymakers and ethicists alike.

In this workshop, you will be creating a Use Case which uses AI to summarise a Car Insurance Claim. The scenario will allow you to experience the following capabilities in watsonx.governance:

- Creation of a Use Case including automated workflows
- Performing Risk and Compliance Assessments
- Gaining approvals for the sign off to Development
- Development of a prompt template
- Test your prompt, try different parameters and models.
- Evaluate your prompt against gold standard summaries.
- Track your prompt through its lifecycle within the AI use case. This includes version control.
- Promote and evaluate your prompt template in a pre-production deployment space.
- Managing the prompt in production.

Introduction to your lab scenario

In this lab, your client is a large insurance company that is attempting to infuse AI into their business, while complying with best practices and regulations. Using the watsonx.governance platform, you will demonstrate how, with a single solution, they can govern, monitor, and document both traditional predictive models and new generative AI models.

Phase 1: Plan


We define a Use Case as a collection of information to solve a business problem, in this case an Insurance Claim Summarisation. In practice, an organization would identify a business need for an AI model, and create a use case to track that effort.

The watsonx.governance solution allows organizations to group and track their models based on use cases, or issues that models are attempting to solve. Each use case stores and organizes data and lifecycle information for candidate models in planning, design, development, and production phases.

In the planning phase, you will use pre-configured workflows to populate the necessary information about the Use Case itself.


Create an AI Use Case

1. Please select the hamburger menu on the top left.




.

2. Select **AI use cases** from the menu.




3. Create a new AI use case. At this point you will be redirected to the watsonx.governance console however once the steps are completed the Use Case will also be visible in your list of AI Use Cases screen. This is an example of the bi-directional flow between the Evaluation/Monitoring and AI Factsheet modules and the Governance Console




4. Fill out the Use Case using the following parameters:
- Give it a **Name** such as “Insurance claim summarisation – (*insert your team name*)”
 - Add yourself as the **Owner**
 - **Primary Business Entity** can be set to any listed
 - Click **Save**
5. Next, fill out details which are applicable to the Use Case itself, e.g.:
- **Description** set “GenAI Experience Workshop - Car Insurance Claim Use Case”
 - *Any others?*

NOTE: If you receive the following error:



Then from Setting disable **“System Admin Mode”**



Note: you may also get an error indicating someone else has updated the Use Case/Object. If you do, please refresh the screen.

- Whilst in Edit mode  set the **Technical Owner** to YOUR USER and hit  at the top right of the screen.

IBM watsonx | Governance console

Insurance claim summarisation

Proposed

Low

CancelSave

TaskActivityAdminSecurity Performance Monitoring

ModifiedRequired

General

NameInsurance claim summarisation

Use Case TypeAIDescriptionGenAI Experience Workshop - Car Insurance Claim Use Case

Ownerisilbmgroup001 isilbmgroup001 isilbmgroup001@ibm.com

PurposeGenAI Workshop

Stakeholder Departments

Legal

Engineering

CTO

Information Security

Ethics

Mentel Blink

Technical Owner

isilbmgroup001 isilbmgroup001 isilbmgroup001@ibm.com

Uses Generative AIDisabled

Externally FacingDisabled

Target Implementation Datemm/dd/yyyy

StageUse Case Data Gathering (Data gathering)

Due Date2/4/2025

Tags

Use Case Data Gathering

All Key Items (5)

Purpose

Risk Level

Use Case Type

Stakeholder Departments


Technical Owner

8

7. The Use case should look like this:

The screenshot shows the IBM Watsonx Governance console interface for a 'Use Case' titled 'Insurance claim summarisation'. The status is 'Proposed' and the risk level is 'Low'. The console is divided into several sections: 'General' (Name, Description, Stakeholder Departments, Third Party Link), 'Use Case Details' (Use Case Type, Owner, Technical Owner, Purpose, Status, Risk Level), and 'Risk' (Risk Level, Risk Identification Completion Date, Risk Assessment Completion Date). A right-hand sidebar provides additional details for 'Use Case Data Gathering', including a due date of 2/4/2025, a list of tags, and a section for 'All Key Items (5)' to validate, which includes Purpose, Risk Level, Use Case Type, Stakeholder Departments, and Technical Owner.

8. Now that you have successfully created the Use Case you can **Submit it for Initial Approval** by navigating to the **Actions** button at the Top right of the screen:



The Action button only works when the required information is captured. This is all driven by customisable Workflows which are set up to match an organisations business process'. This is outside the scope of the labs however if you wish to see this, please ask the table facilitator to demonstrate.

9. On the pop-up window select **Continue**:

Submit for initial approval

Are you sure you want to perform this action?

Continue

Continue and close tab

10. Once the Use case is Submitted for Approval, the Use Case gets updated and under the **Risk** section further down the screen, an **AI Risk Identification Assessment** has been automatically created:

IBM watsonx | Governance console

Insurance d... Use Cases + New Use Case

Insurance claim summarisation ☆ ^

Task Activity Admin Security Performance Monitoring

* Modified Required *

Use Case Details ⓘ

Uses Generative AI Externally Facing Target Implementation Date

Proposed Solution Additional Details Data Gathering Completion Date

Risk ⓘ

Risk Level Low Risk Identification Completion Date Risk Assessment Completion Date

Risk Identification Assessments

Search Add New

Name	Description	Progress (%)	Tags
☐ AI Risk Identification (Insurance claim summarisation)	AI use case risk identification assessment (2025-01-30)		

Use Case Risks All Risks

Search Copy from Library

Name	Description	Inherent Risk Rating	Residual Risk Rating	Status	Tags
------	-------------	----------------------	----------------------	--------	------

Stage Initial Approval (Awaiting use case approval)

Due Date 2/4/2025

Tags

No tags have been added yet.

Initial Approval ⓘ

Please review the initial details related to the use case as captured by the Use Case Owner.

Use the Actions button to Return to owner, reject the use case, or more

Select an action to validate

All Key Items (3)

Purpose


Risk Level

Use Case Type

Risk Assessment

As part of the planning phase, we need to perform a risk assessment of our AI Use Case. The following steps show how an organisation can use a pre-build risk assessment to identify and manage the AI Risks associated with their Use Case.

1. Select the **AI Risk Assessment** by scrolling down to **relationships** and selecting Questionnaire Assessment → AI Risk Identification and Launch the questionnaire.



2. Go ahead by reading and answering the questions. For the purposes of our exercise, it is more important to understand the nature of the **Questions** themselves, therefore feel free to provide answers such as “**ABC**” & “**No**” :

1. Select the **Assessment**:

The screenshot shows the IBM Watsonx Governance console interface. The top navigation bar includes 'Insurance cl...', 'Use Cases', and '+ New Use Case'. The left sidebar shows 'Questions completed' at 0/19 and 'Sections' with 'Main Section' expanded, containing 'AI Use Case Risk Identification'. The main content area is titled 'Main Section 0/19' and 'AI Use Case Risk Identification 0/19'. It contains an 'Introduction' section with a text area and a 'Clear' button. Below this is a 'Problem statement' section with a text area and a 'Clear' button. The bottom section is 'Expected users' with a text area and a 'Clear' button. At the bottom right, there are buttons for 'Save draft', 'Action', and 'Submit and Close'.

2. Once all **Questions** are answered, go ahead and select **Submit and Close** under the **Action** button:

A close-up of the 'Action' button dropdown menu. The menu is open, showing three options: 'Save draft', 'Action', and 'Submit and Close'. The 'Action' button is highlighted in blue.

3. From the top, **Click** on the tab to navigate back to the **Insurance Claim Summarization Use Case**:

The screenshot shows the IBM Watsonx Governance console interface. The top navigation bar includes 'Insurance cl...' and 'AI Risk Ident...'. The left sidebar shows 'Questions completed' at 20/20 and 'Sections' with 'Main Section' expanded, containing 'AI Use Case Risk Identification'. The main content area is titled 'Questionnaire Assessment' and 'AI Risk Identification (In...'. It contains a 'Have you already selected a model for implementing this use case?' section with radio buttons for 'Yes' and 'No'. Below this is a 'Clear' button. At the bottom right, there are buttons for 'Save draft', 'Action', and 'Submit and Close'.

Note that the **Status** from **Proposed** has progressed to **Awaiting Use Case Approval**.


The screenshot shows the IBM Watsonx Governance console interface. At the top, there's a navigation bar with 'Insurance cl...' and 'Use Cases' tabs. Below this, the 'Insurance claim summarisation' Use Case is selected. The status is 'Awaiting Use Case Approval' (highlighted with a red box) and the risk level is 'Low'. The main panel is divided into sections: General, Use Case Details, and Risk. The General section includes fields for Name, Description, Stakeholder Departments, and Third Party Link. The Use Case Details section includes fields for Use Generative AI, Proposed Solution, and Target Implementation Date. The Risk section includes fields for Risk Level, Risk Identification Completion Date, and Risk Assessment Completion Date. On the right, there's a sidebar with 'Initial Approval' details, including a stage, due date, and a list of key items (Purpose, Risk Level, Use Case Type).

4. Additionally, further down in the **Risk** section, we can see one (or many depending on your answers) Use Case Risk identified.


The risk(s) was populated in the Use Case based on our answers on the Questionnaire in the previous step. The risks are configured as a Response Action to the question and come from the AI Risk Atlas, which is an IBM Research produced set of risks associated with Gen AI and Traditional ML. It can be found here: [AI risk atlas - IBM Documentation](#)

The screenshot shows the 'Risk' section of the IBM Watsonx Governance console. It contains two tables. The first table, 'Risk Identification Assessments', has columns for Name, Description, Progress (%), and Tags. It lists one assessment: 'AI Risk Identification (Insurance claim summarisation)' with a progress of 100%. The second table, 'Use Case Risks', has columns for Name, Description, Inherent Risk Rating, Residual Risk Rating, Status, and Tags. It lists one risk: 'Non-disclosure (Copy of MOD_0000000_RIS_0000038)' with an inherent risk rating of 'Not Determined', a residual risk rating of 'Not Determined', and a status of 'Awaiting Assessment'.

5. Next we will assess the risk in the context of our Use Case. **Open** the Use Case Risk by selecting it:




6. Go ahead and select **Start Model Risk Assessment** under the **Action** button:



At this stage, the person responsible for the risk assessment (e.g. a Risk Owner) would assess the risk based on controls that exist in the environment. You can see there are options to associate any mitigation controls with the risk itself. This is outside the scope of this lab.

7. Once the **Risk Assessment** is initiated proceed by flagging it as **Ready for Assessment** under the **Action** button:



8. Switch to Edit mode  and provide **Low** ratings for **Inherent Impact**, **Inherent Likelihood**, **Residual Impact**, **Residual Likelihood**, as well as a placeholder “**ABC**” for **Migration Strategy**.

The risk level is a subjective rating and calculates an automatic overall risk rating for the specific risk.

9. Finally, **Save** and proceed as **Assessment Complete**. Feel free to add a Workflow Comment with this.

10. Once the Risk Assessment is completed, navigate back to the Use Case.

You will see the Risk(s) associated with the Use Case and the fact it is “Accepted”. You will also see a series of dashboards and risk ratings which can be configured to show the current risk profile of your Use Case. This is a useful visualisation.

It is now the responsibility of the Model Owner to provide a Risk Rating commensurate with the outcome of the Risk Assessment process. This is subjective based on the Risks visible in the Use Case.

11. To do this, scroll down to the **Risk** section, switch to Edit mode  and set the Risk Level to **Low**:

The screenshot displays the IBM Watsonx Governance console interface. The main section is titled "Insurance claim summarisation" and contains a form with several input fields and a dropdown menu for "Risk Level". The "Risk Level" dropdown is currently open, showing options: High, Medium, Low, Not Available, and No response. The "Low" option is selected. Below the form, there is a "Progress (%)" bar and a "Tags" section. The right sidebar shows the "Use Case Data Gathering" stage with a due date of 2/4/2025 and a list of key items: Purpose, Risk Level, Use Case Type, Stakeholder Departments, and Technical Owner.


While this risk assessment comes out of the box, it is simple to create your own Risk Assessments based on your own set of organisational risks. This is outside the scope of this lab.

EU AI Act applicability Assessment

The EU AI Act is a regulation affecting AI Use Cases. It categorises Use Cases depending on the purpose of them from Prohibited, High, Medium and Out of Scope. It is important to understand the applicability of your AI Use Case to the regulation, as there are regulatory fines for non-compliance. If the result was prohibited, you would immediately know that this would never be a production use case hence saving yourself time developing this Use Case.

This pre-configured assessment helps an organisation know under which category of the act their AI Use Case falls under.

1. Complete the **Applicability Assessment** under **relationships** and selecting Questionnaire Assessment → Applicability Assessment and Launch the questionnaire.



2. Complete and submit the questionnaire assessment.

Insurance claim summarization – TestUser... Completion required No Due date Save draft Submit and Close

View all questions

Questions completed 0/3 Copy answers

Sections Applicability Assessment Scope and Prohibited AI Systems

Applicability Assessment 0/3

Scope and Prohibited AI Systems 0/3

1.1.1. Organization Classification

How would you classify your organization?

Review each statement and check all classifications that apply. *

Refer to [Article 3. Definitions](#) for the definition of Provider, AI system, Deployer, Importer, Distributor, Authorized Representative and all other applicable definitions.

☐ A provider placing on the market or putting into service AI systems or placing on the market general-purpose AI models in the Union, irrespective of whether your organization is established or located within the Union or in a third-country

☐ A deployer of AI systems that have their place of establishment or who are located within the Union

☐ A provider and deployer of AI systems that have their place of establishment or who are located in a third country, where the output produced by the system is used in the Union

☐ An importer of an AI system

☐ A distributor of an AI system

☐ A product manufacturer placing on the market or putting into service an AI system together with the product and under your own name or trademark together with your product

☐ An authorized representatives of a provider, which are not established in the Union

☐ Affected persons located in the Union

☐ None of the above

Comment Activity

1.1.1.1. Excluded AI Systems

The AI System has been developed or will be used for the following purposes:

Review each statement and check all answers that apply. *

☐ Placed on the market, put into service, or used with or without modification exclusively for military, defense or national security purposes, regardless of the type of entity

Submit for stakeholder review

Assuming this wasn't prohibited, we are now ready to submit this for stakeholder review. The list of stakeholders indicate the different disciplines involved in the sign off of an AI Use Case. This is a final step before the AI Use Case is signed-off for development and ensures that the requisite stakeholders in an organisation have the opportunity to review all the info relating to the Use Case. If they are not happy with this they can reject it or request further information.

For this step, you will be taking on the persona of someone responsible from the legal team signing off on the Use Case.

1. Switch back to the **Use Case** and proceed with **Submit for stakeholder review**:

The screenshot displays the IBM Watsonx Governance console interface. The main section shows the 'Insurance claim summarisation' Use Case. The status is 'Awaiting Use Case Approval' and the risk level is 'Low'. The 'Actions' menu is highlighted, showing options: 'Reject use case', 'Submit for stakeholder review', and 'Return to owner'. The 'Submit for stakeholder review' option is the target for this step.

General

Name: Insurance claim summarisation

Use Case Type: AI

Status: Awaiting Use Case Approval

Description: GenAI Experience Workshop - Car Insurance Claim Use Case

Owner: isibmgrou001.isibmgrou001

Purpose: GenAI Workshop

Stakeholder Departments: Legal

Technical Owner: isibmgrou001.isibmgrou001

Third Party Link: <https://dataplatfom.cloud.ibm.com/aigov/modelinventory/inventories/62288f7e-ca4e-4acb-8997-2d4168a37900/aiusecase/f0965380-6c11-44db-a3a8-2381746d1926>

Use Case Details

Uses Generative AI: Externally Facing

Target Implementation Date:

Proposed Solution: Additional Details

Data Gathering Completion Date:

Risk

Risk Level: Low

Risk Identification Completion Date:

Risk Assessment Completion Date:

Initial Approval

Please review the initial details related to the use case as captured by the Use Case Owner.

Use the Actions button to Return to owner, reject the use case, or more

Select an action to validate

All Key Items (3)

- Purpose
- Risk Level
- Use Case Type

2. Once submitted, it is ready to be reviewed by one of the stakeholders, as it can be seen in the **Use Case Approvals** section:

The screenshot displays the IBM Watson Governance console interface. The top navigation bar includes tabs for 'Insurance cl...', 'Use Cases', and 'Non-disclos...'. The main content area is titled 'Insurance claim summarisation' and features a 'Task' tab. Below this, there's a section for 'Use Case Approvals' with a table listing review items. One item is visible: 'Insurance claim summarisation-Review-00003' with a 'Legal' stakeholder department, 'Awaiting Approval' status, and 'Not Determined' risk rating. To the right, the 'Stakeholder Review' section provides instructions on how to approve use cases for development, accompanied by a progress bar and a list of key items for attention.

Use Case Approvals

Stakeholder Approval Completion Date

Use Case Reviews

Name	Stakeholder Departments	Approval Date	Approval Status	Risk Rating	Review Comment	Tags
Insurance claim summarisation-Review-00003 Reference: Innovation Studio	Legal		Awaiting Approval	Not Determined		

Related Models

All Models | Prompts and Tunes | Associated Foundation Models | Deployments | Model Links | Model Groups

Search

Name	Description	Model Owner	Model Class	Model Status	Tags
No results					

Other Relationships

Stakeholder Review

Any relevant stakeholder departments that were selected in the data gathering stage have been sent a Use Case Review task to obtain their approval for this use case. Once all of the approvals are received, this use case can be approved for development.

Approve for development

1 item requires attention.

All Key Items (4)


- Purpose
- Risk Level
- Use Case Type

Key Items for this Action


Logic: 1 OR 2

- 1: Use Case Reviews
- 2: Use Case Reviews

3. Open the **Use Case Review** by selecting it:



4. Set the **Reviewer** as your current username, for **Review Comment** type “**Approved**”, set the **Risk Rating** as Low and press **Save**:



5. Once Saved, **Approve Use Case** under the Action Button:



6. Return to the Use Case. You can now **Approve for Development**:



7. The Use Case has progressed through the Plan phase and had the necessary approvals from key stakeholders including Risk Management and Legal. This Use Case is now “Approved for Development”.

The screenshot displays the IBM Watsonx Governance console for a Use Case named 'Insurance claim summarisation'. The status is 'Approved for Development' (highlighted with a red box) and the risk level is 'Low'. The interface includes tabs for Task, Activity, Admin, and Security Performance Monitoring. The main content area shows details for the Use Case, including Name, Description, Stakeholder Departments (Legal), and Third Party Link. A right-hand sidebar contains a 'Tags' section and a 'Use Case general view' section with a progress bar and key items (Purpose, Risk Level, Use Case Type).

8. Scroll down to the bottom of the Use Case to see that there has been a relationship diagram developed that ties the Use Case to the various objects recently developed.

The screenshot displays the IBM Watsonx Governance console for a Use Case named 'Insurance claim summarisation - sammclear'. The 'Relationships' section at the bottom shows a diagram connecting the Use Case to four entities: Business Entities, Model Groups, Questionnaire Assess..., and Risks. The diagram uses colored dots to represent different relationship types: Primary Parent (blue), Parent (purple), and Child (black).

9.

Well done, you have successfully completed Phase 1 of the AI Use Case lifecycle, the Planning phase. You have identified:

- the purpose


- technical owners
 - the Risk Level
 - whether it is compliant to the EU AI Act - stakeholder approval.
- Please pause here before we explain the next Phase.

Phase 2: Design


Create a Prompt Template

Let's begin creating our first prompt using Watsonx.AI. Watsonx.AI is IBM's AI studio that allows you to build, tune and deploy AI models.

You can click on the third party link within the Use Case to navigate to the Watsonx.AI platform which will show you the Use Case. Separately use this link: <https://ibm.biz/BdGSj2>



From IBM watsonx.AI, use the hamburger menu to navigate to your Team's project e.g. Team 1.




Projects are spaces to collaborate with others and build data and AI assets. You will see multiple options along the top which include information about the project, the assets associated with the project, any current jobs that are running and to allow the management of that asset.


Create a New Project, name the project “Insurance Claims – YOURNAME”

Select the Cloud Object Storage-td for the storage service.


Head over to the **Assets** tab of the project home page. Once there, you will be prompted to associate a watsonx.ai Runtime service, select **Associate Service**.




Select the “watsonx.ai Runtime-td” service and click **Associate**.




Head back over to the **Assets** tab of the project home page. Once there, click on **New Asset**. Note that your project will already include assets for a later part of the lab.



Now, if you scroll down on the menu that appears, you should see a card that says **Chat and build prompts with foundation models**.



You will be taken into the Watsonx.AI prompt lab. If there is an option to take a tour, skip it. You will be met with the prompt lab. This screen has numerous options on how you can configure your prompt to suit your needs.




Right above that, you see the option to select “Chat”, “**Structured**” or “Freeform”.

Structured prompts have easily identifiable sections to place examples, instructions, and prompt variables (which we will discuss later). Freeform prompts have a looser and more flexible structure. With Chat mode, you can chat with the foundation model to see how the model handles dialog or question-answering tasks.

This lab will walk through utilising the **Structured** prompt functionality, however this is possible to do with freeform as well, so choose whichever option is the most comfortable for you!

You will notice that there is a box labelled **Instruction** which is where we will be placing the bulk of our prompt to be evaluated later.



You will also see an icon with a graph at the top right , which is the button to run an evaluation that

we will discuss later. You will be able to see the status of your work at the top right as well, indicating whether your prompt session work is Saved or Unsaved.

Select the “llama-3-3-70b-instruct” model from the menu on the top right.

Inside of the Instruction box, enter your first prompt that we will be evaluating throughout this lab. For now, enter the prompt shown in the screen below. This is the instruction given to the model.

Copy & Paste from in between the quotation marks:

“You are an insurance agent tasked to assess insurance claims. Summarise the following insurance claim input. Focus on the car and the damage. Make the summary at least 3 sentences long.”

Projects / Insurance claims - TEAM NAME / Prompt Lab

Unsaved New prompt + AI guardrails off

Chat Structured Freeform

AI Model: llama-3-3-70b-instruct

Hint: This model works better when you provide at least 1 example.

Set up

Instruction (optional)

You are an insurance agent tasked to assess insurance claims. Summarize the following insurance claim input. Focus on the car and the damage. Make the summary at least 3 sentences long.

Examples (optional)

Input:	Output:
Enter your example input here.	Enter your desired output.

Add example +

Prompt Variables

Now, we need to ensure that we are able to test our prompt against validation and test data. To do so, we need to ensure we have prompt variables set up for our instructions.

2783672 - Innovation Stu... Dallas TC

New prompt + AI guardrails off

AI Model: llama-3-3-70b-instruct

Prompt variables

Once you have opened up this sidebar, you will see two empty boxes pop up, with **Variable** and **Default Value** able to be filled out. Use "input" as the variable name and leave the default value empty.

Prompt variables ⓘ

Variable	Default value
input	default value

New variable +

Preview ⓘ

By scrolling down your page, you will notice an area on the screen titled **Try**. Inside of this, you want to put a variable titled "{input}". This variable will be replaced by input insurance claims.

Projects / Insurance claims - TEAM NAME / Prompt Lab

Unsaved ⓘ New prompt + AI guardrails off ⓘ

Chat Structured Freeform

Model: llama-3-3-70b-instruct

Hint: This model works better when you provide at least 1 example.

Set up ^

Instruction (optional) ⓘ
You are an insurance agent tasked to assess insurance claims. Summarize the following insurance claim input. Focus on the car and the damage. Make the summary at least 3 sentences long.

Examples (optional) ⓘ

Input:	Output:
Enter your example input here.	Enter your desired output.

Add example +

Try ^

Test your prompt ⓘ

Input:	Output:
{input}	Generated output appears here.

New test +

Clear output

Generate →

Prompt variables ⓘ

Variable	Default value
input	default value


New variable +

Preview ⓘ

Model Parameters

Next to where you clicked prompt variables, you should see an option that says **Model parameters**. Click on that next. We will not be making any serious changes here just yet but mostly giving a tour of this section so you can make modifications later.

Most importantly, when you select the button highlighted below (top right), you will see a toggle for **Greedy** or **Sampling**.




Selecting greedy will allow your model results to be reproduceable across experiments, while selection sampling will give you more variability.

For our task, we will be choosing the sampling option for this base prompt.

A screenshot of the 'Model parameters' configuration panel. The panel has a title 'Model parameters' and a sub-section 'Decoding'. Under 'Decoding', there is a toggle switch for 'Greedy' (which is currently turned on) and 'Sampling' (which is currently selected). Below this are sliders for 'Temperature' (set to 0.7), 'Top P (nucleus sampling)' (set to 1), and 'Top K' (set to 50). There is a text input field for 'Random seed'. Below that is a slider for 'Repetition penalty' (set to 1). At the bottom, there is a section for 'Stopping criteria' with a sub-section 'Stop sequences' containing a text input field and a '+' button. Below that are two input fields for 'Min tokens' (set to 0) and 'Max tokens' (set to 200). A note at the bottom says 'Enter up to 6 sequences to stop output after the minimum number of tokens is reached.'

Now that you have chosen sampling, you will see several new options pop up that are specific to the sampling decoding method. **Temperature** is the most important factor here, where setting a higher temperature will give higher variation and more unique outputs. Hover over the info icons to gain more information on **Top P** and **Top K** methods as well.

Lastly, scroll all the way down the **Model parameters** sidebar to find the **Stopping criteria** as well as **Token limits**. For LLMs, tokens ~ words, so we want to set that higher (e.g. 200) to get more detail in our summaries of the insurance claims.



The screenshot shows a sidebar titled "Stopping criteria" with a help icon. It contains a "Stop sequences" section with a text input field and a "+" button. Below this are "Min tokens" and "Max tokens" sections, each with a text input field. The "Min tokens" field is set to "0" and the "Max tokens" field is set to "200". At the bottom, there is a note: "Enter up to 6 sequences to stop output after the minimum number of tokens is reached." and a "Reset to default" button with a circular arrow icon.

Saving Prompt

Once you are all finished, click the **Save** icon dropdown in the top right, and then click **Save as**. This will take you to the following screen, where you will select your asset type, prompt template name, and task.

Make sure to select the **Prompt template** option, then fill out a name in a similar fashion shown in the picture. From there, select **Summarization** for your task, and make sure to select **View in project after saving**.


You have now successfully designed your first prompt template.

Phase 3: Development


Evaluating your prompt template

The goal of this part of the lab is to iterate on our first prompt that we have created, and work to achieve better results so that we can deploy our prompt and eventually bring it to production.

Now you should see your finished prompt template on the screen when you are taken back to the **Assets** page. Verify that you have completed the previous steps correctly by clicking the 3 dots to the right of the prompt and choose **Evaluate**. We are going to be testing out our new prompt!



You should now see a screen pop up with an evaluate button in the middle. Select that button.



Now you will see a screen asking you which dimensions to evaluate. At this point, we are only able to evaluate GenAI quality (model health is enabled by default). By clicking on the arrow next to the Generative AI Quality, you can see which metrics are being calculated and what the definition of those are. It is also possible to switch metrics on/off via the Advanced settings feature.

The screenshot shows the 'Evaluate prompt template' interface. On the left, a sidebar contains four steps: 'Select dimensions' (active), 'Select test data', 'Map variables', and 'Review and evaluate'. The main area is titled 'Select dimensions to evaluate' with a subtitle 'These dimensions are based on the prompt template task type. [Learn more](#)'. A table lists two dimensions: 'Generative AI Quality' (checked) and 'Model health' (unchecked). A description for 'Generative AI Quality' is provided. An 'Advanced settings' link is in the top right. At the bottom are 'Cancel', 'Back', and 'Next' buttons.

☒ Dimension	Description
☒ Generative AI Quality	The Generative AI Quality monitor calculates a variety of metrics based on prompt template task type. Some metrics compare model output to the reference output you provide. Other metrics analyze model input and output and do not require reference output.
☐ Model health	The model health monitor provides an overview and helps to understand how your model deployment is performing with the incoming transactions. This monitor evaluates and computes metrics such as scoring requests count, records count, latency and throughput, etc..

Now it's time to choose our test data. For this tutorial, we will be using the **Insurance claim summarisation validation data.csv** from your lab's Data folder.

Choose Select from Project and locate the file.

The screenshot shows the 'Evaluate prompt template' interface. On the left, a sidebar contains four steps: 'Select dimensions' (selected), 'Select test data' (active), 'Map variables', and 'Review and evaluate'. The main area displays a file upload prompt with a file icon labeled 'Insurance claim summari...data.csv'. Below the icon, it says 'Drop a file here or browse for a file to upload'. To the right, there is explanatory text: 'Add a CSV file that includes input and expected output (ground-truth). Test data for this deployment can optionally include model output to remove the need for additional model transactions. Maximum size is 8 MB. Maximum number of records is 1000. Minimum number of records is 10.' Below this text are two buttons: 'Browse' and 'Select from project'. At the bottom, there are three buttons: 'Cancel', 'Back', and 'Next'.

Now it's time to map the input variables to the input variable that we created earlier. Additionally, I chose **Input** to be mapped to the “Insurance_Claim” column in my file. This will ensure that my prompt is focused on generating the correct data.

Lastly, **Reference output** is the column of the data that gives the expected output from our LLM where metrics will be evaluated against. This is how we can measure how effectively we have structured our prompts.

I chose the reference output to be mapped to ‘Summary’ – this is the column of data that gives the expected output to the LLM.

The screenshot shows the 'Evaluate prompt template' interface at the 'Map variables' step. The sidebar on the left shows 'Map variables' as the active step. The main area is titled 'Map prompt variables to columns' and includes a sub-header 'For each prompt variable, select the associated column, Learn more'. It contains two sections: 'Field separation' with a dropdown menu set to 'Comma (,)', and 'Input' with a dropdown menu set to 'Insurance_Claim'. Below these is a 'Reference output' section with a dropdown menu set to 'Summary'. At the bottom, there are three buttons: 'Cancel', 'Back', and 'Next'.

After that, you will see a screen to review the information you have entered. Once you verify that you have entered everything correctly, hit the **Evaluate** button. and wait until your prompt has finished being evaluated. This process can take up to a few minutes.


The screenshot shows the 'Evaluate prompt template' interface. On the left, there is a sidebar with four steps: 'Select dimensions', 'Select test data', 'Map variables', and 'Review and evaluate' (which is currently selected). The main area is titled 'Review' and contains the following information:

- Task:** Text summarization
- Test data:** Insurance claim summarization validation data.csv
- Evaluations:** Generative AI Quality, Model health


Below this information, there is a note: 'Note: Evaluation can take a several minutes to complete. You can continue to work on other things while your evaluation is in progress.' At the bottom of the interface, there are three buttons: 'Cancel', 'Back', and 'Evaluate'.

Evaluation Results

Once the evaluation has finished, you should see a screen pop up displaying the results. As you can see here, our evaluation was not very successful, triggering a number of alerts and failing its test. This is to be expected!




Click on the blue arrow → on the right of **Generative AI Quality – Text Summarisation** to access a more detailed view of the evaluation results. Make sure to check both the **Model health** and **Generative AI Quality** tabs.



In this lab's use case, the LLM's task was “Summarisation” which can be measured using a range of quality metrics. We recommend taking time now to scan through this documentation on the [generative AI quality metrics supported by watsonx.governance](#).


AI Factsheet

Up in the top left corner of your screen, choose the **AI Factsheet** tab, which will show us a better breakdown of the results we got. The first thing you will see on the page is an option to track your prompt in an AI use case.



First explore the page by clicking into each of the tabs:

- The **Governance** tab collects basic information such as the name of the AI use case, the description, and the approach name and version data.
- The **Foundation model** tab displays the name of the foundation model, the license ID, and the model publisher.
- The **Prompt template** shows the prompt name, ID, prompt input, and variables, whilst the **Prompt parameters** collect the configuration options for the prompt template, including the decoding method and stopping criteria.
- Finally, the **Evaluation / Unspecified phase** tab displays the data from evaluation, including alerts and metric data from the evaluation. Within the **Test** sub-tab you'll see a screen similar to the one shown below:




Note: The metrics that you see in the lab may differ slightly to what you see above.


You can see that each metric has an upper and lower bounds with the ideal range indicated by the bolder portion at the left side of each metric's measurement. The upside-down triangles ▼ represent the measured value for each metrics.

Now, I'll show you how to modify the thresholds for each metric so that your prompt can have looser requirements to pass testing. In a real-world use case, this should only be done if necessary and makes sense for your use case! Head back over to the evaluation screen, where you will see a blue settings button to the right of **Generative AI Quality – Text Summarisation** where we can

adjust our thresholds.




Go ahead and click on that button, and you will be greeted with the following screen:



Here we are going to adjust the **Readability** metric. This metric is higher if your model's generated responses to your prompt are easy to interpret, and lower if they are confusing to read. Scroll down the page and click the edit button to modify it.



Lastly, go ahead and change the readability slightly, say to 55, which puts the requirement into the "Fairly difficult to read" range. After, that, you can go ahead and click save to finalise your changes. Just a reminder here that these metrics should not be altered too much in a real-world use case!




Govern your prompt template


The next section of this lab is to govern your prompt template.

In order to do this, you will need to first associate the project that you created earlier (most likely the Sandbox environment, unless you created one) to the relevant AI Use Case.

To do so, click into the [Track in AI use case](#) button within the **Governance** tab:



Then, select [Go to AI use cases](#) to open up a list of AI use cases.




Select the Use Case that you created earlier, and you will be met with this screen:

The screenshot displays the 'Overview' tab of an AI use case management interface. The main content area is divided into several sections: 'Welcome to your new AI use case' with 'Discover next steps', 'Set your AI use case status' (with a dropdown menu), and 'Get started with samples'. Below these is the 'General information' section, which includes fields for 'Name' (Insurance claim summarisation - TEAM NAME), 'Description' (GenAI Experience Workshop - Car Insurance Claim Use Case), 'Owner' (Murali Krishnan Sankaran), 'Status' (Proposed (System, Jan 27, 2025)), 'Risk level' (None), 'Inventory' (WatsonX Workshop), and 'Tags'. A 'Details' section follows, with a 'Purpose' field and a 'Supporting documentation' field. At the bottom is the 'Associated workspaces' section, which contains a button labeled 'Associate workspace +'. A right-hand sidebar titled 'About this AI use case' provides a summary of the use case details, including name, description, owner, status, risk level, inventory, and tags.

Scroll down to the **Associated workspaces** section. In this lab, the relevant phase is the '**Develop**'; the prompt requires engineering to develop and further refine it. Click the button shown below to [Associate workspace +](#).

The screenshot shows the 'Associated workspaces' section of the interface. It includes a heading 'Associated workspaces' and a sub-heading 'Associate your AI use case with the workspaces in order to organise them under the same business problem.' Below this, there are two main sections: 'Develop' and 'Validate'. The 'Develop' section is highlighted with a blue bar on the left and contains the text: 'In this phase your development team will experiment with different approaches of ML development or prompt engineering and test them in projects.' A red rectangle highlights the 'Associate workspace +' button. The 'Validate' section is partially visible at the bottom.

You will see the following screen and will need to select the relevant **Project** (most likely your Sandbox environment) by ticking the box on the right.




Tracking Prompt Template in a Use Case

Now that you have associated your project to your AI use case, you can track your prompt template in your AI use case. The idea behind tracking prompt template in a Use Case is to save AI Engineers time and give transparency into the work they're doing. You can automatically save prompt templates and use versioning to easily identify the different prompts and where they are in the lifecycle.

Select the navigation menu in the top right again, and under **Projects**, click into **View all projects** to navigate to your project. Click into this and find the prompt you made earlier. If you struggle to find your way to your project, your facilitator will be able to help.

If you click the three dots on the right, you will be able to see that the prompt needs to be tracked – click on the Track in AI usecase button.



Select the **Default approach**, select to automatically create a **New asset record** and make it **Experimental**.

Insurance Claim Summary Prompt

Track in AI use case

Track an asset to collect details about the asset in factsheets as part of a governance strategy.

Define approach

Define asset record

Assign version

Review

Define approach

Use case: **Insurance claim summarisation – TEAM NAME**

An approach defines one path for solving the goal of the use case. Ex. an approach might be a variation on a machine learning model, or a challenger model. Each approach can include multiple versions.

New approach +

Default approach

A default approach for track-
ing your AI assets.

Cancel

Back

Next

Insurance Claim Summary Prompt

Track in AI use case

Track an asset to collect details about the asset in factsheets as part of a governance strategy.

Define approach

Define asset record

Assign version

Review

Define asset record

Approach: **Default approach** | Use case: **Insurance claim summarisation – TEAM NAME**
Associate the trained model with an existing model record or create a new model record in OpenPages. This will sync the tracked model facts between Model inventory and OpenPages.

Asset Record

Existing asset record

+ New asset record



Create record automatically

A new asset record will be created in OpenPages when model tracking is initiated.

Cancel

Back

Next

Insurance Claim Summary Prompt

Track in AI use case

Track an asset to collect details about the asset in factsheets as part of a governance strategy.

Define approach

Define asset record

Assign version

Review

Assign version

Approach: **Default approach** | Use case: **Insurance claim summarisation – TEAM NAME**
Choose the starting point for this approach.

Experimental

Use this as a starting point if your model is just starting in development and its input and output structure will likely change in the near future.

0.0.1

Stable

Use this as a starting point if your model is in a production state and you won't expect any major changes in its input and output structure soon.

1.0.0

Custom

Define your own starting version if you already tracked this model in a versioning context before.

Version number

0.0.1

Comment (optional)

Provide context for this change. Comments are saved in the model submit to govern.

Cancel

Back

Next

Review your choices and start Tracking the Asset

Insurance Claim Summary Prompt

Track in AI use case

Track an asset to collect details about the asset in factsheets as part of a governance strategy.

Define approach

Define asset record

Assign version

Review

Review

Important

Make sure your prompt template is stable before you track it in an AI use case. Tracking a prompt template locks it so it is no longer editable. Complete edits and evaluations before enabling tracking. You can untrack the prompt template to edit it, but you will lose the tracking history.

Insurance claim summarisation – TEAM NAME

Approach

Version

Default approach

0.0.1

A default approach for tracking your AI assets.

00000000-0000-0000-0000-000000000000

Cancel

Back

Track asset

Projects / Insurance claims – TEAM NAME / Insurance Claim Summary Prompt

Open in Prompt Lab

AI Factsheet

Evaluate

Governance

Foundation model

Prompt template

Prompt parameters

Development

Insurance claims – T...

Test results

Additional details

Attachments

Governance

Export report

Insurance claim summarisation – TEAM NAME

Proposed | System | WatsonX Workshop | None | c2428369-0ef4-4f06-996d-04b345363efc

Description

GenAI Experience Workshop - Car Insurance Claim Use Case

Approach

Version

Default approach

0.0.1

A default approach for tracking your AI assets.

00000000-0000-0000-0000-000000000000

Lifecycle

01 Develop

02 Validate

03 Operate

Untrack

Foundation model

Using the hamburger menu, navigate to your **AI Use Case**, switch to the **Lifecycle** tab see your prompt template lifecycle being tracked. If it doesn't appear immediately refresh your screen.

46

The screenshot displays the IBM Data Platform interface for an AI use case titled "Insurance claim summarisation - TEAM NAME". The interface is divided into three tabs: "Overview", "Lifecycle", and "Access", with "Lifecycle" currently selected. At the top right, there is an "Export report" button and a help icon.

Below the tabs, there are three informational cards:

- Welcome to your new AI use case**: Includes a "Discover next steps" link.
- Organize team work with approaches**: Explains that different paths can solve a business problem and that an approach is used to organize assets for one solution.
- Manage AI assets with versions**: States that each approach starts a new version series for identifying associated assets and that a particular state of assets can be preserved in one approach.

Below these cards, there is a toggle for "Show deleted assets" and a "New approach +" button.

The main section is titled "Default approach" with a "Latest: 0.0.1 | 1 version" indicator. It describes the approach as "A default approach for tracking your AI assets." and shows a lifecycle flow with three stages:

- Development**: Experiment with different approaches of ML development or prompt engineering and test them in projects.
- Validation**: Evaluate models and prompt templates with validation data.
- Operation**: Evaluate and monitor models and prompt templates in production.

Below the lifecycle flow, there is a version selector "v0.0.1". Under this version, there are three boxes representing the state of assets:

- Development**: Contains a box for "Insurance claims - TEAM NAME" with a sub-item "Insurance Claim Summary Prompt" marked with a red dot.
- Validation**: Contains the text "No models promoted to a pre-production space."
- Operation**: Contains the text "No models promoted to a production space."


This is where you can view any changes to the prompt template as we continue with the lab.

Update your prompt template


We will now make additional changes to the prompt template, which we'll do by creating a new version of the original prompt.

For that, use the main menu to navigate to your project. We'll create a new prompt that is a copy of the first one. Do that by opening your template from the **Assets** tab and use **Save as** to give it a new name, such as "Insurance Claim Summary Prompt v2". That opens it for editing again.

Change the model to "**granite-13b-instruct-v2**" to introduce a change in this new version of the prompt template!




Evaluate your prompt template by following the same steps as before ([Evaluating your prompt](#)).




Once you have completed the evaluation of this new prompt template, you can track it in your AI use case. Go to your project by selecting the project name in the path at the top.

Track it in your AI use case by selecting **Minor Change (version 0.1.0)**.



Check that it appears in the **Lifecycle** tab of your AI use case.




You will see that the two prompts are now tracked separately within the same AI use case, making it easy to track different prompts, models or parameters for a single use case/application.


Evaluation Studio

The purpose of the Evaluation Studio is to look at different prompts and identify which is the strongest performing. For the lab, you will compare the two prompt templates you've created. Note if you didn't get this far, then there are existing prompts you can use. The following lab explains how simple it is to perform an evaluation using the Evaluation Studio.


Navigate back to your project and select "New Asset".




In the new UI, search for Evaluate and Compare prompts.




Leave the name and description as it. Select the task type as Summarisation




Select the two different prompts you created earlier. In this example they are labelled Insurance Claims Summarisation Attempt 1 and 2. Click Next.




Review the different metrics which you are evaluating against. Again these fall under the categories General AI Quality and Model Health and can be configured. .




On the Test Data screen, please select the Insurance Claims Summarisation test data csv from the project. This is under the project file




Similarly to before, the Input is “Insurance Claim” and the Output is “Summary”



Review the evaluation before running it, this will take a few minutes.



Spend some time going through the results, identifying how each attempt performs differently across several criteria.



Once done, navigate back to the project nad you will see these results exist as an evaluation result.

Phase 4: Deployment

If we consider that you've iterated the prompts and found the right combination of model, parameters, model variables until you're happy. You'd now look to promote that asset into production. When you put an asset in production, there are a different set of Risks associated with it for example Drift. This next part of lab will show you how to promote the asset into production and the features available plus turning on the evaluation for drift.

Deployment Space

A prompt template is only useful if it is accessible by an application; this prompt template needs to be made available in a consumable way. What's more, the deployment process needs to be completed in an orderly manner. For example, you cannot allow everyone to put prompt templates/models into production. For this purpose, the watsonx.ai environment includes deployment spaces.

Anyone can create their own personal deployment space and promote prompt templates/ models to it. However, in an enterprise, designated administrators should be the only ones able to promote a model to a production deployment space.

Go back to your Use Case and scroll down to Workspaces, click to add a Workspace in Validation:

Go back to your Use Case and scroll down to Workspaces, click to add a Workspace in Operate:

Overview Lifecycle Access

Model risk governance details

Additional Details ⓘ
Data not available from OpenPages

Associated workspaces

Associate your AI use case with the workspaces in order to organise them **under the same business problem**.

Development

Team 5

Validation

In this phase your validators can evaluate models in pre-production deployment spaces and prompt templates in projects.

Associate workspaces +

Operation

In this phase your development team will evaluate models in pre-production deployment spaces.

Associate workspaces +

AI use case

Workspaces for Develop

Workspaces for Validate

Workspaces for Operate

Click on the button shown below to create a [New space](#).



You will see the window shown on the next page.

Fill in the details as specified below:


- Name the new deployment space something appropriate
- Select the storage service and machine learning service from the options that are available under the respective dropdown menus on the right side
- Select the Deployment stage as 'Validation'
- Select the Storage as "Cloud Object Storage-td"
- Associate the "watsonx.ai Runtime-td service"
- Click **Create**

Once you have created this deployment space, press [Save](#) to associate it to the workspace **Validation**.



Next, find your second prompt version (the one created during the Develop phase) within **Projects** and promote it to this new workspace. Select the Deployment space you created earlier.

We then need to deploy the prompt, click on the 3 buttons next to your asset and choose deploy:



Now repeat the same steps in **Deployment Space** for a Workspace in **Operate**:


Fill in the details as specified below:

- Name the new deployment space something appropriate
- Select the storage service and machine learning service from the options that are available under the respective dropdown menus on the right side
- Select the Deployment stage as 'Production'

- Press [Create](#).
-


The screenshot shows the 'Create a deployment space' form. The title is 'Create a deployment space' with a subtitle 'Use a space to collect assets in one place to create, run, and manage deployments'. The form is divided into two main sections: 'Define details' and 'Select services'. In 'Define details', there are fields for 'Name' (filled with 'Deployment Space - Production'), 'Description (Optional)' (filled with 'Deployment space description'), 'Deployment stage' (a dropdown menu with 'Production' selected), 'Deployment space tags (optional)' (a dropdown menu with 'Find or create tags' selected), and 'Advanced Settings' (a dropdown menu). In 'Select services', there are two dropdown menus: 'Select storage service' (filled with 'Cloud Object Storage-pj') and 'Select watsonx.ai Runtime service (optional)' (filled with 'Watson Machine Learning-a1'). Below these is an 'Upload space assets (optional)' section with a text box that says 'Populate your space with assets exported from a project or space to a .zip file. You can add more assets after the space is created.' and a button that says 'Drop .zip file here or browse your files to upload'. At the bottom right of the form are two buttons: 'Cancel' and 'Create'.

Once you have created this deployment space, press [Save](#) to associate it to the workspace **Operate**.




Next, find your second prompt version (the one created during the Develop phase) within **Projects** and promote it to this new workspace. Select the Deployment space you created earlier.

We then need to deploy the prompt, click on the 3 buttons next to your asset and choose deploy:




To confirm the steps worked, you should see the prompt template in the *Production deployment space* on the far-right panel, with the *Pending Evaluation* label




Well done, you have successfully deployed your first asset into production.

Managing Assets in production

Click on the recently created production deployment space (the one that is shown to be pending evaluation), and click [Open in Space](#).




The first thing to talk through is the API reference. As discussed, there is no point in developing a model if it isn't easy to access. With the deployment space, you can see that there are public and private API endpoints to allow an application to call the model. Making it even simpler, there are code snippets available in multiple languages to allow this to be baked into the application code. Thus, making our production asset readily accessible.




Now, we can test it to make sure it is functioning as expected.


Navigate to the **Test** tab within this **Deployment Space** to ensure the function is working as expected.



Within this tab, click on the [Generate](#) button to make sure the prompt works.




Review the schemas of the payload and the feedback data. The datasets (csv) files need to follow these schemas when we upload them.



We are now ready to prime the prompt deployment with baseline data.

Key	Value
<i>Generated text</i>	Car A hit Car B from behind, totalling Car B.
<i>Generated token count</i>	14
<i>Input token count</i>	45
<i>Stop reason</i>	End of sequence token encountered
<i>Seed</i>	0

Go to the **Evaluations** tab and click the [Activate](#) button.



You should notice that now you have activated the monitoring, you will automatically have a new dimension called **Drift v2** as an additional option. As shown in the screenshot below, the description gives details as to how it measures distribution shifts between training and payload data – which we must now upload.

Evaluate prompt template

Choose the evaluation dimensions and select the test data. [Learn more](#)

Select dimensions

Review and evaluate

Select dimensions to evaluate

These dimensions are based on the prompt template task type. [Learn more](#)

Column names in payload data must match the prompt variable names in the prompt template. Feedback data must include an additional label column for reference output. Provide the name of the label column in the field below. Note that the Drift v2 evaluation requires 100+ records to establish a baseline.

Label column

Dimension

Description

☒

Drift v2

Computes distribution shifts between training and payload data.

☐

Embedding drift

Embedding drift is defined as the percentage of records where the runtime data was detected as an outlier as compared to the baseline data. This is an opt-in experience for the user. Once enabled, the user needs to provide embeddings along with the data for this metric to be computed.

☐

Output drift

Output drift is defined as the change in distribution for the model output between the baseline data and the runtime data. For classification problems, this is computed on the individual class probabilities. For regression problems, this is computed on the model prediction. For the generative AI models, this is computed on the log probability of the final generated text.

☐

Input metadata drift

Input metadata drift is defined as the change in distribution for the model input metadata between the baseline data and the runtime data. If the tag, 'field_type', is present and set to 'meta', it will denote drift in the input metadata column in 'field_name' tag. Otherwise, this metric is at subscription level.

☐

Output metadata drift

Output metadata drift is defined as the change in distribution for the model output metadata between the baseline data and the runtime data. If the tag, 'field_type', is present and set to 'meta', it will denote drift in the output metadata column in 'field_name' tag. Otherwise, this metric is at subscription level.

☒

Generative AI Quality

The Generative AI Quality monitor calculates a variety of metrics based on prompt template task type. Some metrics compare model output to the reference output you provide. Other metrics analyze model input and output and do not require reference output.

Cancel

Back

Next

Once you have walked through the steps, click [Activate](#) to activate monitoring of the prompt in production. This asset is now being monitored in production on a scheduled basis, metrics captured will be sent to the Governance Console and tracked within the Use Case you created earlier.

Note: The Prompt Results that you see in the lab may differ slightly to what you see above.

IBM watsonx

2783672 - Innovation Stu... Dallas 10

Deployments / Deployment Space - Production / Insurance Claim Summary Prompt v2 /

Insurance Summarisation - Prod Deployment

Deployed Online

API reference

Test

Evaluations

AI Factsheet

Last evaluation: --

Deployment details

Test details

Model health

0 Tests run

Tests passed 0 Tests failed 0

Records 0

Token count 0

Latency (record) 0

ms Median record late

Total input token count

Total output token count

Generative AI Quality - Text summarization

Not evaluated


Awaiting model transactions to evaluate. Evaluation will run when the minimum number of transactions is met.

Drift Monitoring (Optional – feel free to progress to the [Governance Console](#) section)

To demonstrate drift we need to establish a baseline. Consider this as the ground truth for the data. That is to say at time of production the data is up to date and accurate. Over time of course, Drift will be introduced as the data becomes outdated and no longer relevant. The calculation of this works off the baseline data.

Select Upload Payload data and then upload the file called:

Insurance claim summarisation drift payload 0 - baseline.csv



Evaluate

Import test data

Description

For summarization, generation, and question answering tasks, watsonx.governance evaluates generative AI quality (reference free metrics), drift, and model health using logged scoring requests as received by the model. Scoring requests are logged using the payload logging endpoint or file upload. [Learn more.](#)

For summarization, generation, and question answering tasks, watsonx.governance optionally evaluates production models for generative AI quality (reference based metrics) using labeled test data. Labeled test data is provided using the feedback endpoint or file upload. [Learn more.](#)

Upload payload data

[View payload schema](#)

Upload feedback data

[View feedback schema](#)

Total records

Payload records	1
Feedback records	0

Close **Evaluate now**

Note: We are not uploading any feedback data at this point.

Once that has been successfully uploaded, Click on the [Evaluate now](#) button.

The **Evaluation** will fail due to a violation of the **Coverage** Metric. We could of course amend the threshold of the metric for our **Evaluation** to Pass, however for the purposes of our workshop will accept it as it is.

If you now click on the [Configure monitors](#) button in the [Actions](#) dropdown menu, you will see **Drift v2** enabled.

As shown in the screenshot below, you will have a baseline of 100 records. You can edit the threshold values here as well which would be important in a real-world scenario, to adjust them to an appropriate level given the use case and prompt.

For now, we'll keep the default values.

Insurance Summarisation - Prod Deployment

Model info

Model details

Evaluations

Fairness

Quality

Drift v2

Explainability

Generative AI Quality

Model health

Drift v2

Description

The Drift monitor checks if your deployments are up-to-date and behaving consistently. Model input/output data is analyzed in relation to the training/baseline data.

Compute the drift archive

Compute option

Compute in Watson OpenScale

Baseline data

Number of records

100

Sample size

Minimum sample size

10

Maximum sample size

1,000

Drift thresholds

Upper thresholds

Output drift

0.05

Input metadata drift


0.05

The system will now monitor invocations of the prompt and detect any potential drift. We will simulate an example of this by uploading the following datasets.


- Insurance claim summarisation drift payload 1.csv
- Insurance claim summarisation drift payload 2.csv
- Insurance claim summarisation drift payload 3.csv
- Insurance claim summarisation drift payload 4.csv

Upload all of the csv files as payload data, but be sure to let some time pass between each run, to be sure that the evaluation shows a meaningful timeline (and drift).

Once you have run these evaluations, you will see a sub-tab under **Evaluations** to show a diagram for the drift of input and output. Click on the  blue arrow button in the corner of the Drift v2 section to get a more detailed view.



Make sure you have selected the right time settings to see the drift change over time. If you have any issues, ask for assistant from your facilitators.



There are only three criteria which go into the drift calculations: *output*, *output metadata drift* and *input metadata drift*.


These metrics do not offer any information regarding the *quality* of the prompt – In order to measure that, we need to upload *feedback data*.

Navigate back to the **Evaluations** tab and once again click on the [Evaluate now](#) button within the [Actions](#) dropdown menu. Upload the following file:

- Insurance claim summarisation feedback.csv


When this run is finished, you will see a panel with **Generative AI Quality - Text summarization** metrics on the lower right of the screen, next to the Drift v2 section. If the option allows, click on the  blue arrow button in the top right corner to get a more detailed view of the metrics collected.

Drift v2			Generative AI Quality - Text summarization		
Alerts triggered			Alerts triggered		
Alerts			Alerts		
3			1		
Metric	Score	Violation	Metric	Score	Violation
Output drift	0.20	0.15	Coverage	0.55	0.15
Input metadata drift	0.20	0.15	Density	0.13	--
Output metadata drift	0.20	0.15	Compression	1.98	--
			Repetitiveness	0.09	none
			Abstractness	0.12	none
			Readability	89.42	none
110 records evaluated			110 records evaluated		



Finally, we are ready to advance the status of the use case, to identify it as a model in **Operation**.

The final tab, AI Factsheet, captures all of the facts about the AI Asset. It is also possible to view information such as the serving name on the right hand side of the screen.



The prompt is deployed in production and is being monitored for both drift and quality. As seen back in the lifecycle screen.

The screenshot displays the IBM Watsonx interface for an AI use case titled "Insurance claim summarisation - TEAM NAME". The interface is divided into three tabs: Overview, Lifecycle (selected), and Access. The Lifecycle tab shows a "Default approach" for tracking AI assets, with a timeline from Development to Operation. The Development phase includes a "v0.1.0" version of the "Insurance claims - TEAM NAME" approach, which contains an "Insurance Claim Summary Prompt v2". The Validation phase shows the "Insurance Claim Summary Prompt v2" being evaluated in the "Deployment Space - Validation". The Operation phase shows the "Insurance Claim Summary Prompt v2" being evaluated in the "Deployment Space - Production". The v0.0.1 version shows no models promoted to pre-production or production spaces.

IBM watsonx

AI use cases / Insurance claim summarisation - TEAM NAME

Export report

Overview Lifecycle Access

Welcome to your new AI use case

Discover next steps

Organize team work with approaches

Manage AI assets with versions

Show deleted assets

New approach +

Default approach

Latest: 0.1.0 | 2 versions

A default approach for tracking your AI assets.

Development

Experiment with different approaches of ML development or prompt engineering and test them in projects.

Validation

Evaluate models and prompt templates with validation data.

Operation

Evaluate and monitor models and prompt templates in production.

v0.1.0

Insurance claims - TEAM NAME

Insurance Claim Summary Prompt v2

Deployment Space - Validation

Insurance Claim Summary Prompt v2

Insurance Summarisation - Deployment

Deployment Space - Production

Insurance Claim Summary Prompt v2

Insurance Summarisation - Prod Deployment

v0.0.1

Insurance claims - TEAM NAME

No models promoted to a pre-production space.

No models promoted to a production space.


Governance Console

For the last part of the lab, you will see how the work carried out throughout this lab was automatically captured in the Use Case on the Governance Console.

Navigate back to the Governance Console either from the Use Case view or this link

<https://ibm.biz/BdGmZc>

Click back into your teams' Use Case and scroll down to the Performance Monitoring section:



Scroll down to the Relationships screen. You can see the different objects that are now tied to the Use Case. To drill down further into the models themselves and their performance, click on the name and additional objects will appear.

The screenshot displays the IBM Watsonx Governance console interface. The top navigation bar shows the 'Insurance claim summarisation - sammclear' use case. The main content area features a 'Relationships' diagram with a central node 'Insurance claim summ...' connected to several other nodes: 'Business Entities', 'Model Groups', 'Model Links', 'Models', 'Questionnaire Assess...', and 'Risks'. The right-hand sidebar contains a 'Tags' section with the message 'No tags have been added yet.' and a 'Use Case general view' section with a progress bar and a list of key items: 'Purpose', 'Risk Level', and 'Use Case Type'. Below the main content is a 'Related Models' section with a table of models.

Name	Description	Model Owner	Model Class	Model Status	Tags
flan-t5-xl-3b Library > MRG > Foundation Models	flan-t5-xl (3B) is a 3 billion parameter model based on the Flan-T5 family. It is a pretrained T5 - an encoder-decoder model pre-trained on a mixture of supervised / unsupervised tasks converted into a text-to-text format, and fine-tuned on the Fine-tuned Language Net (FLAN) with instructions more		Foundation Model	Proposed	
sam 2 Watsonx Innovation Studio			Prompt-based	Proposed	
sams Insurance prompt Watsonx Innovation Studio			Prompt-based	Proposed	

Given the evaluation results, we may want to retrospective actions such as Issue Management or change the Risk level to Medium.

IBM watsonx | Governance console

Use Case: Insurance claim summarisation - TEAM NAME

Status: Proposed

Risk Level: Medium

Task: Activity Admin Security Performance Monitoring

* Modified Required *

General

Name: Insurance claim summarisation - TEAM NAME

Status: Proposed

Description: GenAI Experience Workshop - Car Insurance Claim Use Case

Owner: isibmgrou001 isibmgrou001

Purpose:

Third Party Link: <https://dataplatfom.cloud.ibm.com/aigov/modelinventory/inventories/622887e-ca4e-4acb-8997-2d41d8a37900/aiusecase/c24283e9-0ef4-4f06-99ed-04b345363efc>

Use Case Details

Additional Details:

Risk

* Risk Level: Medium

Regulatory Information

Use Case Data Gathering

Stage: Use Case Data Gathering (Data gathering)

Due Date: 2/1/2025


Tags: No tags have been added yet.

Please capture all relevant information to this AI use case proposal and then submit using the Action button.

All Key Items (2)

- Purpose
- Risk Level

Congratulations! You have completed the Workshop.
You have built, deployed and managed an AI Use Case
tackling the 3 pillars of AI Governance.





Optional steps: Navigating the UI of the Governance Console


Get Started with your watsonx.governance Workshop

1. Open **Chrome** or **Firefox** and navigate to the following Link: <https://ibm.biz/BdGmZc>
2. **Logon** with the following Credentials:
 - a. Username:
 - i. islibmgroup001@ibm.com
 - ii. islibmgroup002@ibm.com
 - iii. islibmgroup003@ibm.com
 - iv. islibmgroup004@ibm.com
 - v. islibmgroup005@ibm.com
 - b. Password:
 - i. 13thFeb2025\$WorkShop

If an OTP is required, please ask your table facilitator to retrieve this.

Hint: make sure all spaces are removed from start and end of UN & Password

watsonx.governance User Interface



After login, you will be displayed the Homepage  tab, and in it we will see four tabs.


- **HOME** tab. The dashboard displays panels with content that is personalized for each user, intended to be relevant information for the user. *You may move and arrange them, as per your preferences. Later we will see how to create a new panel.*

The next three Task related tabs, are displayed tasks that I am notified during the workflow lifecycle of the solution components (*Click on each tab to see the content*):

- **My Tasks:** activities that I have been assigned. It is action that I must perform.
- **Subscription Tasks:** activities that I want to know that happen, however I do not have to do anything.
- **Oversight Tasks:** I have full visibility on the whole lifecycle of an element (i.e., a Model Use case). I do not have any task assigned, just visibility.

Changing Colour Theme

1. Themes are used to customize the colours that are displayed in OpenPages (watsonx.governance).
2. Users can apply a theme by using the  > Change Theme task.



3. Using the Drop-down list, change your Theme to Carbon Gray 90.




If you do not like the contrast, you may choose any other.

Themes can be configured by the client to allow a corporate colouring of the solution.

The Primary Menu


1. On the top left corner, click on the Primary Menu  (also known as the “Hamburger menu”)




2. The Primary menu lists the categories of object types that you can access in IBM OpenPages® (watsonx.governance).

Expand and collapse the categories to see the object types. Seeing more or less objects depends on your profile.

3. When you click an object type, for example, **Use Cases** under the section Inventory, a list of objects opens.



4. Click on , expand **Inventory**, click on **Use Cases**.



Name	Purpose	Description	Owner	Status	Risk Level	Tags
Agency Based LGD Estimation High Oaks Bank > North America > Corporate Banking		Uses internal and external recovery data, adjusted for macro-economic impact. Uses statistical regression	Bob Eldridge	Approved for Development	Low	
Banking book HTM corporate bond - income High Oaks Bank > Europe > Corporate Banking		ALM based income forecast for the HTM portfolio, initially for the CCAR 2013 stress-test. Vendor solution using conditional scenarios and core ALM system. Linked to fair-value and impairment models.	Bob Eldridge	Approved for Development	Medium	
Black model for IR derivatives High Oaks Bank > North America > Investment Banking		Black Linear-Nonlinear model on IR process	Bob Eldridge	Approved for Development	Medium	
CCAR Stress Test High Oaks Bank > North America > Investment Banking		Stress tests are submitted according to macroeconomic scenarios provided by the Fed	Bob Eldridge	Approved for Development	High	
Commodity Options VaR High Oaks Bank > North America > Investment Banking		Stochastic VaR at 99.97%. Pricing model Mapping: Asian Commodities: Cost; American Commodity Option: Risk Metrics; European Option: Black-Scholes-Merton	Bob Eldridge	Approved for Development	High	
Conditional scenarios High Oaks Bank > North America > Investment Banking		Conditional scenario creation from main risk drivers using linear regression	Bob Eldridge	Approved for Development	High	
Conversational AI High Oaks Bank > Corporate		AI system that is capable of understanding, processing and responding to human language in a conversational way.	Bob Eldridge	Approved for Development	High	
Credit Risk High Oaks Bank		Model to evaluate credit risk of loan applications.	Bob Eldridge	Approved for Development	High	
Customer Attrition High Oaks Bank > Europe > Retail Banking		A system to predict which customers will cancel their services	Bob Eldridge	Approved for Development	Medium	
Customer Chatbot High Oaks Bank > Corporate		A chatbot to assist customers with finding the right service and answer questions.	Bob Eldridge	Approved for Development	Medium	
CVA - WWR adjustment		Adjustment on CVA price due to Wrong Way Risk, in the portfolio correlation between exposures and probability of default	Bob Eldridge	Approved for Development	Medium	

The list of Use Cases Display. This type of view is called Grid View. When you select an object type from the Primary menu, the **Home**, or a **Task View**, a Grid View is displayed.


There are several actions that might be taken from a grid view.

Search Use Cases


1. Whilst in the Use Cases Tab, place the cursor on the Search box.



2. Start Typing **“Governance”**. You are presented with a Tag Suggestion



3. Select the Tag **“Governance”**.


A screenshot of a header area titled "Use Cases (2)". To the right of the title is a grey pill-shaped button with the text "Governance" and a small "X" icon to its right, indicating an active filter.

Notice the filter indicator displayed.

4. Place the cursor on the Search box once again.

A screenshot of a search bar. On the left is a funnel icon. To its right is a magnifying glass icon followed by the text "Search".

5. Type "**Insurance**". This is a Text Filter.


A screenshot of a header area titled "Use Cases (1)". To the right of the title is a grey pill-shaped button with the text "Governance" and a small "X" icon. Below the header is a search bar containing the text "Insurance". Below the search bar is a table with two columns: "Name" and "Purpose". The first row of the table has a checkbox, the text "Wx.demonstration Test", and the text "Watsonx Innovation Studio".

Notice that both filters have been applied.

6. Remove the filter display **Governance**.
Notice that now we have only a text search based on the word "insurance".

These filters are **Ad hoc** and cannot be reused, as you must define them each time. Ad hoc filters are private filters that are not saved and only used in one session. Use an ad hoc filter if you do not frequently access the search criteria.

7. Click on the funnel symbol on the left to open the filter side bar.



Notice that we have a section named “Public Filters” and “My Filters” (private filters)

An administrator defines public filters. Users can customize private filters for their own use.

Use and customize private filters to find objects based on criteria on field values on objects and, optionally, parent objects. Users can customize private filters for their own use.


Only one public or private filter can be applied at a time.

You cannot simultaneously use a public filter and a private filter.

Let us create a filter that we may reuse.

Create Search Filters

1. Click on **New Filter +**.



Filter


Business Entity [Select +](#)

Public Filters

My Use Cases 0

My Filters [New Filter +](#)

2. Provide the following **Name**: Governance & Legal Compliance




New Filter

* Name

Governance & Legal Compliance

3. Click on **New Condition +**.



* Name

Governance & Legal Compliance

Base Object Criteria ^

Conditions [New Condition +](#)

ID	Condition
No Conditions	

4. Select **Tag &** from the drop-down list select Governance & Legal Compliance.

New Condition

* Select a condition type.

☐ Field ☒ Tag

* Tags *

2 x selected items

☐ Accuracy

☐ Data laws

☐ Explainability

☐ Fairness

☒ Governance

☐ Harmful code generation

5. Click **Save**.

Notice that now appears your filter under “My filters”.

×

Filter

Business Entity

Select +

Public Filters

My Use Cases0

My Filters

New Filter +

Governance & Legal Compliance 4



6. Your Filter is now displayed by the object type text.

Use Cases (4) Governance & Legal Compliance ×

- Remove the filter and return to the homepage tab.




Configuring your home page: create a new panel

Now we are going to create a new panel, as we want to focus temporarily on Governance & Legal Compliance Use cases. We are going to design a panel that uses our just created private filter.


1. Click on the **Configure** icon on the right, above the tabs.



2. Select **New Panel +**.



3. From the drop-down list select **Chart**.



4. Accordingly provide the following parameters:

Label: Governance & Legal Compliance by Risk

Object Type: Use Case

Filter: Governance & Legal Compliance

Chart Type: Pie (you may also want to choose any other, except Gantt, to see how it looks like)

Chart Data Field: Risk Level

Method Type: Count

Home /

New Panel

* Label *

Governance & Legal Compliance by Risk

* Object Type *

Use Case

* Filter

Governance & Legal Compliance

* Chart Type *

Pie

* Chart Data Field *

Risk Level

* Method Type *

Count

Sort Direction

Field definition order

Cancel Done

5. Click **Done & Save**.

A new panel will appear in your dashboard.



6. Click on the categories in the legend to **hide/unhide** them.



7. Select any section on the chart to drill down on a Risk Level category and you will be taken to a grid view displaying your filter name and the filter based on the risk level you have clicked on.



Use Cases (1) Governance & Legal Compliance Risk Level: High

Search

☐ Active Only

New

☐	Name	Purpose	Description	Owner	Status	Risk Level	Tags
☐	Credit Risk High Oaks Bank		Model to evaluate credit risk of loan applications.	Bob Eldridge	Approved for Development	High	Legal compliance

8. If you wish to remove your panel from the homepage, click on the **Configure** icon on the right, above the tabs.



9. In the displayed list of panels, you would look for the one you would want to remove, hover the mouse over it and you would remove it by clicking on the minus sign.

×

Configure dashboard

Customize the panels on the current tab and manage all tabs.

Panels

Tabs

Home tab

New Panel

+

Change Management

Deployments

Governance & Legal Compliance by Risk

⊖

Delete

Issues by Status

Model Compliance Status

Model Inventory

Model Performance

Model Validation

Models by Department

Models by Provider

Do NOT remove the newly created panel as we will use it later on.

Navigation from an object

1. From the Homepage click on the “High Risk” level area in the chart of the panel you have just created.



2. Do a quick inspection of the record by clicking on the **Open Side Panel** on the right. 

A screenshot of a web application showing the details of a 'Credit Risk' Use Case. The form is titled 'Use Case Credit Risk' and includes a star icon for favoriting and a search icon. Below the title, there are fields for 'Name' (Credit Risk), 'Owner' (Bob Eldridge), 'Status' (Approved for Development), 'Risk Level' (High), 'Purpose', 'Description' (Model to evaluate credit risk of loan applications), and 'Third Party Link'. The form is divided into sections: 'General' and 'Associations'. The 'Associations' section has tabs for 'Models' and 'Deployments'. At the bottom, there are 'Cancel' and 'Action' buttons.

3. Mark this Use Case as Favorite (it will make it appear in your “My Favorites” panel in your homepage), by clicking on the star icon to flag/unflag the record as Favorite.

A screenshot of the 'Credit Risk' Use Case details form, similar to the one above. A tooltip is visible over the star icon, stating: 'Add to favorites. They are listed in the favorites dashboard panel.' The tooltip is a dark gray box with white text. The form itself is partially obscured by the tooltip.

4. Select the Use Case to open it in a new tab, as we want to investigate further.

The screenshot displays the IBM Watsonx Governance console interface. At the top, the breadcrumb navigation shows 'Use Cases' > 'Credit Risk'. The main header area includes the 'Use Case' title 'Credit Risk', a star icon, and a status badge 'Approved for Development'. To the right, there is a 'Risk Level' badge 'High' and an 'Action' button. Below the header, a tab bar shows 'Task', 'Activity', 'Admin', and 'Security Performance Monitoring'. The main content area is divided into sections: 'General' and 'Associations'. The 'General' section contains fields for 'Name' (Credit Risk), 'Owner' (Bob Eldridge), 'Status' (Approved for Development), 'Purpose', and 'Third Party Link'. The 'Associations' section has tabs for 'Models' and 'Deployments'. Below the 'Models' tab, there is a search bar and a table listing associated models.

Name	Description	Model Owner	Model Status	Third Party Link	Tags
☐ MOD-00001 High Oaks Bank	Credit Risk	Mark Owens	Pre Implementation Review		

5. Scroll down to the section **Related Models**. The Tab **All Models** displays the Models related to this Use Case.

This screenshot is a zoomed-in view of the 'Associations' section from the previous image. It shows the 'Models' tab selected, with a search bar and a table of associated models.


Name	Description	Model Owner	Model Status	Third Party Link	Tags
☐ MOD-00001 High Oaks Bank	Credit Risk	Mark Owens	Pre Implementation Review		

6. Scroll down to **Other Relationships**. The relationships tree view is displayed.




7. Open the tree diagram in a new tab by clicking on the **View fullscreen** icon 

8. Following the picture below, click on the circles to expand the content. Click on the individual item to drill down on it. Notice the blue area that explains the relationships between the objects.




Audit Trail

1. Click back into the tab for the Credit Risk Use Case



2. Switch from the **Task** tab to **Activity** tab.



3. You may see all the changes made to this element. The details cannot be edited.

The screenshot displays the Audit Trail table for the Credit Risk Use Case. The table has five columns: Activity type, Old value, New value, Modified by, and Date and time. It contains two rows of activity data.

Activity type	Old value	New value	Modified by	Date and time
Tag - Added		Legal compliance	Tasos Christodoulou	Jan 24, 2025 6:38:44 AM EST
Model - Added		MOD-00001	System Administrator	Dec 5, 2024 1:48:43 AM EST


Additional UI elements visible in the screenshot include the 'IBM watsonx | Governance console' header, navigation tabs for 'Use Cases', 'Credit Risk', and 'Activity' (selected), and a status bar at the top right showing 'Status: Approved for Development' and 'Risk Level: High'.

Promoting your prompt to Validation

N.B. This stage is skipped to save time. In real life, enterprises would you validation as a testing environment.


In between development and production there is a validation step. In this next section, we will walk through how to promote and validate your prompt through the stages of the model lifecycle.

To start, navigate back to the **Overview** tab.



Since we now want to promote the second prompt template to the validate phase, we will create a new deployment space and associate it to the **Validation** stage.

Click on the [Associate workspace +](#) button under the Validate space.



We will create a new **Deployment Space** for validation, to promote your prompt into.

Give the **deployment space** a unique name as appropriate to this **validation stage** (i.e Deployment Space - Validation), so we can distinguish from other **deployment space** we will be creating.

Select **Testing** as the deployment stage, as well as the default storage service and machine learning service as you did in the earlier section, before clicking [Create](#).

Create a deployment space

Use a space to collect assets in one place to create, run, and manage deployments

Define details

Name
Deployment Space - Validation

Description (Optional)
Deployment space description

Deployment stage ⓘ
Development

Deployment space tags (optional) ⓘ
Find or create tags

Advanced Settings

Select services

Select storage service ⓘ
Cloud Object Storage-94

Select Watson AI Runtime service (optional) ⓘ
Watson Machine Learning-a1

Cancel Create

Associate this new Deployment Space under the *Space* section then click [Save](#).

Once this is completed, navigate back to your Project, and press the menu for the most recently created prompt template, and click Promote to space

Name	Last modified	
 Insurance Claim Summary Prompt Prompt template	45 minutes ago Modified by you	I
 Insurance Claim Summary Prompt v2 ⓘ Prompt template	2 hours ago Modified by you	I

Evaluate


Go to AI factsheet

Untrack


Promote to space

Delete

Select the newly created deployment space under the **Target space** dropdown menu, and once completed click the [Promote](#) button.




Navigate to the **Lifecycle** tab in your AI Use Case and note that it should appear under the Validate stage as opposed to the Develop stage of the AI use case lifecycle.




The prompt has not been deployed yet – so this is what we will be doing in the next section.

Deploying the prompt


Use the navigation menu to move into the **Deployments** section.




Click into the relevant validation space (i.e. the space you created in the previous section) within the **Spaces** tab.




Move into the on the **Assets** tab, and click the three dots to the right of the most recent prompt template as shown:




Finally, click the option to Deploy the prompt.



Within the deployment wizard, give an appropriate name (i.e. Insurance Summarisation - Deployment) before clicking [Create](#).




If you navigate back to the AI use case, you will see the prompt as Pending Evaluation.



This is the stage at which an appropriate stakeholder would evaluate and start monitoring the prompt – which we will be doing in the next and final section.

Activate prompt monitoring (Evaluation)

Click on the prompt that is pending evaluation, and move into the **Evaluations** tab to be taken to the screen shown below:



Click on the [Evaluate](#) button in the centre of the screen to start another evaluation before checking the results. It may seem redundant to run evaluations with the same prompt and the same data set over and over, but it is important to remember:

- The evaluations would be performed by different people in different roles, e.g. the prompt developer would not be the one evaluating it for use in production.
- In a real-world scenario, the data sets would differ between evaluations.
- There is a strong likelihood that the different models will change the evaluation scores.

After clicking [Evaluate](#), navigate back to the AI use case to confirm that the status has changed from Pending to **Evaluated**

